API usages:

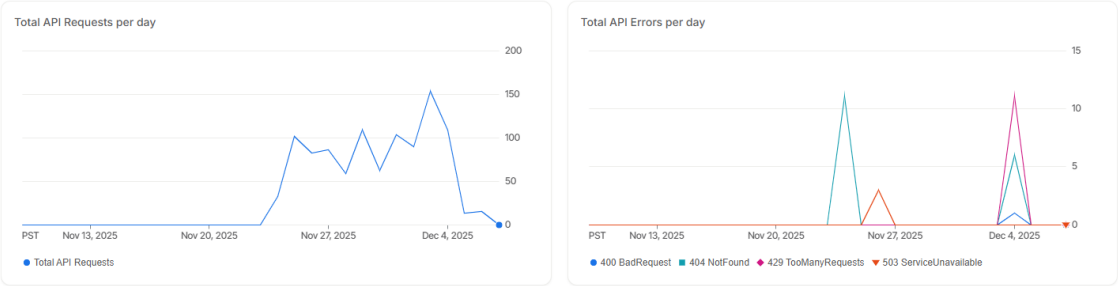
Total API Requests per day
200
150
100
50
0
PST
Nov 13, 2025
Nov 20, 2025
Nov 27, 2025
Dec 4, 2025
Total API Requests
Total API Errors per day
15
10
5
0
PST
Nov 13, 2025
Nov 20, 2025
Nov 27, 2025
Dec 4, 2025
400 BadRequest
404 NotFound
429 TooManyRequests
503 ServiceUnavailable

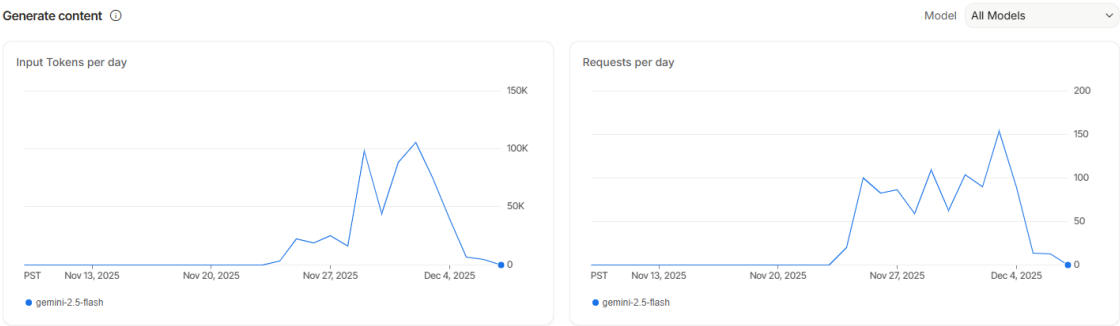
Generate content
Model
All Models
Input Tokens per day
150K
100K
50K
0
PST
Nov 13, 2025
Nov 20, 2025
Nov 27, 2025
Dec 4, 2025
gemini-2.5-flash
Requests per day
200
150
100
50
0
PST
Nov 13, 2025
Nov 20, 2025
Nov 27, 2025
Dec 4, 2025
gemini-2.5-flash

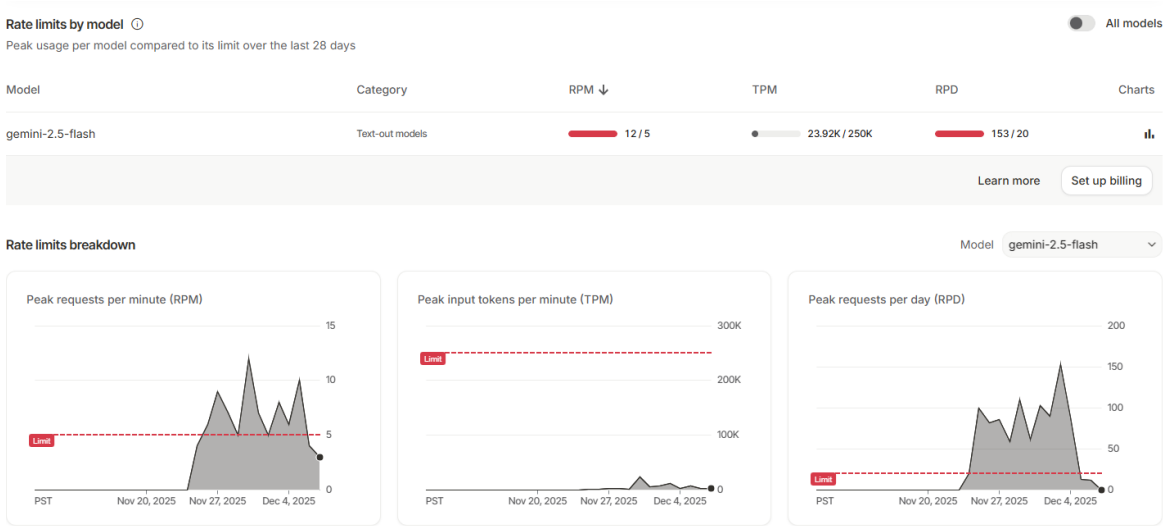
Rate limits by model
Peak usage per model compared to its limit over the last 28 days
All models
Model
Category
RPM
TPM
RPD
Charts
gemini-2.5-flash
Text-out models
12 / 5
23.92K / 250K
153 / 20
Learn more
Set up billing
Rate limits breakdown
Model
gemini-2.5-flash
Peak requests per minute (RPM)
15
10
5
0
Limit
PST
Nov 20, 2025
Nov 27, 2025
Dec 4, 2025
Peak input tokens per minute (TPM)
300K
200K
100K
0
Limit
PST
Nov 20, 2025
Nov 27, 2025
Dec 4, 2025
Peak requests per day (RPD)
200
150
100
50
0
Limit
PST
Nov 20, 2025
Nov 27, 2025
Dec 4, 2025

Gemini 2.5 Flash

`gemini-2.5-flash`

Try it in Google AI Studio

Our first hybrid reasoning model which supports a 1M token context window and has thinking budgets.

Standard Batch

	Free Tier	Paid Tier, per 1M tokens in USD
Input price	Free of charge	$0.30 (text / image / video) $1.00 (audio)
Output price (including thinking tokens)	Free of charge	$2.50
Context caching price	Not available	$0.03 (text / image / video) $0.1 (audio) $1.00 / 1,000,000 tokens per hour (storage price)
Grounding with Google Search	Free of charge, up to 500 RPD (limit shared with Flash-Lite RPD)	1,500 RPD (free, limit shared with Flash-Lite RPD), then $35 / 1,000 grounded prompts
Grounding with Google Maps	500 RPD	1,500 RPD (free), then $25 / 1,000 grounded prompts
Used to improve our products	Yes	No

Why it failed (Usage till now)

- **The Killer Limit:** You have a limit of **20 Requests Per Day (RPD)**.
- **Your Usage:** You are hitting **153 RPD**.
- **The Result:** You are exceeding your daily allowance by **~750%**. This is why your project crashes every day.

Pricing for Gemini 2.5 Flash:

Input: ~$0.30 per 1 million tokens

Output: ~$2.50 per 1 million tokens (Note: Output is pricier because it includes "thinking" tokens, but input is dirt cheap).

Assumption: Your project gets back online and you triple your activity.

- **Daily Usage:** 500,000 Input tokens | 100,000 Output tokens
- **Monthly Math:**
  - Input: 15M tokens $\times$ $0.30 = **$4.50**
  - Output: 3M tokens $\times$ $2.50 = **$7.50**
- **Estimated Monthly Bill: ~$12.00**

Your "200k/day" Request (for current state)

- Input Cost: 200,000 tokens/day × 30 days = 6 Million tokens.
  - $0.30 × 6 = $1.80
- Output Cost: (Assuming ~10% of input size, typical for analysis tasks) = 600k tokens.
  - $2.50 × 0.6 = $1.50
- Total Monthly Cost: $3.30 USD

Tier 2: Heavy Scaling (10x Current Usage)

Assumption: The project becomes very active, handling larger files or more frequent queries.

- Daily Usage: 2 Million Input tokens | 400,000 Output tokens
- Monthly Math:
  - Input: 60M tokens times $0.30 = $18.00
  - Output: 12M tokens times $2.50 = $30.00
- Estimated Monthly Bill: $48.00

Tier 3: "Exponential" Load (50x Current Usage)

Assumption: You are running automated loops constantly or processing massive datasets daily.

- Daily Usage: 10 Million Input tokens | 2 Million Output tokens
- Monthly Math:
  - Input: 300M tokens times $0.30 = $90.00
  - Output: 60M tokens times $2.50 = $150.00
- Estimated Monthly Bill: ~$240.00

Important(To be careful while implementing)

• If your project uses the "Grounding" feature (connecting the model to live Google Search results), you get **1,500 requests per day for free**.

• **After 1,500 requests:** It costs **$35.00 per 1,000 requests**.

• *Risk:* If your code loops and triggers 10,000 search queries, that alone could cost you **$300+** instantly.

Provider	Free Limit	Grounding Cap Type
Google Gemini	**1,500 / day**	**Daily Reset**
Tavily (Popular for Agents)	1,000 / **month**	Monthly Hard Limit
Bing Search API	1,000 / **month**	Monthly Hard Limit
Perplexity API	None (Paid only)	Starts at ~$5/mo
You.com API	Limited Trial	"Basic queries" only